

Problem 2

(a)

Considering the adaptive rejection sampling (ARS) algorithm, we build a set of p points, x_1, \dots, x_p and an upper envelope function that is linear on the log scale withing each interval, i.e. $f(x) = \exp(a_i + b_i x)$ for $x \in (x_{i-1}, x_i)$

We need to normalize this exponentiated upper envelop function in order to have a density of interest that we can draw samples from.

The normalization factor N_f is constant,

$$N_f = \sum_{i=1}^p \int_{x_{i-1}}^{x_i} \exp(a_i + b_i x) dx = \sum_{i=1}^p \frac{e^{a_i}}{b_i} [e^{b_i x_i} - e^{b_i x_{i-1}}].$$

The probability associated with each interval P_i is,

$$P_i = \text{Prob}(x \in (x_{i-1}, x_i)) = \frac{1}{N_f} \int_{x_{i-1}}^{x_i} \exp(a_i + b_i x) dx = \frac{e^{a_i}}{N_f b_i} [e^{b_i x_i} - e^{b_i x_{i-1}}].$$

(b)

We want to draw x from the i th interval using the inverse CDF method.

The CDF function $F(x')$ is,

$$F(x') = \text{Prob}(x \in (x_{i-1}, x')) = \frac{1}{N_f} \int_{x_{i-1}}^{x'} \exp(a_i + b_i x) dx = \frac{e^{a_i}}{N_f b_i} [e^{b_i x'} - e^{b_i x_{i-1}}]$$

We will draw $Z \sim \text{Unif}(0, 1)$ and set $x = F^{-1}(z)$ for i th interval to get $X \sim F$,
 $x = F^{-1}(z) = \frac{1}{b_i} \{\log[N_f b_i z + \exp(a_i + b_i x_{i-1})] - a_i\}$

(c)

The psuedo-code style software that would implement ARS is listed below.

```
# Problem 5.2: ARS [Reference]
# http://www.amsta.leeds.ac.uk/~wally.gilks/adaptive.rejection/web_page/Welcome.html

# calculate the intercept and slope of the line set by succceutive points
lineCoef <- function(x, y) {
  # x in ascending order
  return(data.frame(intercept = a, slope = b))
}

# calculate the intersection points above density curve given the input
# points on the density curve
intersec <- function(x, y) {
  # x and y are the coordinates of the n points
  # get line info for n-1 intersection points
  ab <- lineCoef(x = x, y = y)
  # extrapolate intersection points outside the curve
  qx[i] <- (a[i + 1] - a[i - 1]) / (b[i - 1] - b[i + 1])
  qy[i] <- a[i - 1] + b[i - 1] * qx[i]
  return(cbind(qx, qy))
}

# get the set of points that construct the envelop
envelop <- function(x, y) {
  # x and y are input points on the curve
  Q <- intersec(x, y) # all the upper points
```

```

T <- PQ[order(PQ)] # in order
return(data.frame(T))
}

# function that initializes input to sampling algorithm; The inner and env
# element include the points and lines info for the inner chords and
# envelopes respectively. The region element includes the line info for
# each region that is needed to be calculated
init <- function(x, y) {
  # order x and y in the increasing order of x
  inner <- list(pts = data.frame(x = x, y = y), line = lineCoef(x, y))
  env <- list(pts = envelop(x, y), line = lineCoef(xy.env[1], xy.env[2]))
  reg <- rbind(ab.inner, ab.env)
  return(list(inner = inner, env = env, region = ab.reg))
}

# function to compute the areas of the envelop
area <- function(pts, lb = NULL, rb = NULL) {
  # init from initialization
  # evaluation regions and points
  reg <- pts$region
  x <- pts$env$pts$x
  # calculate the area under the envelope
  ar[i] <- (exp(a + b * x[i + 1]) - exp(a + b * x[i]))/b
  # normalize area and compute cumulatives
  sum.ar <- sum(ar) # raw total area
  ar <- ar/sum.ar # normalized - sum to 1
  cum.ar <- cumsum(ar) # cumulative
  return(list(ar = ar, cum.ar = cum.ar, sum.ar = sum.ar))
}

# Draw sample from envelope
sample.env <- function(pts, lb = NULL, rb = NULL) {
  # calculate the area under the envelope
  ar <- area(pts, lb = lb, rb = rb)
  reg <- pts$region
  x <- pts$env$pts$x
  # line coefficients for that region
  a <- reg[rn, 1]
  b <- reg[rn, 2]
  # sample a point from the envelope using inversion method
  u <- runif(1)
  # invert cumulative df to get sample point
  sx <- (log(u * b * ar$sum.ar + exp(b * x[rn] + a)) - a)/b
  return(sx)
}

# test whether the sample should be accepted or not and store the
# evaluation information of the target function
test.sample <- function(sx, pts, func) {
  x <- pts$env$pts$x
  u <- runif(1)
  if (u < exp(r[1] - r[2])) {

```

```

    # simple acceptance step
    accept <- 1
  } else {
    # evaluation and rejection step
    if (u < exp(func(sx) - r[2])) {
      accept <- 1
    }
    if (!(sx %in% x)) {
      # updating step: update pts
      x <- c(x, sx)
      y <- c(y, sy)
      pts <- init(x, y)
    }
  }
  return(list(accept = accept, pts = pts, xeval = x))
}

# main function to implement the ARS algorithm; n: number of samples
# wanted; func: the log density; xinit: initial values for the X's, at
# least 3; lb: left bound; rb: right bound; The output is a list of 4
# elements; sample: the sample values; xeval: the new x's that cause the
# func to be evaluated; pts: the points and line info for the chords and
# envelop;
ars <- function(n, func, xinit, lb = NULL, rb = NULL) {
  # initialize
  s <- c()
  yinit <- func(xinit)
  pts <- init(xinit, yinit)
  # sample
  while (length(s) < n) {
    sx <- sample.env(pts, lb = lb, rb = rb)
    tsx <- test.sample(sx, pts, func)
    if (tsx$accept)
      s <- c(s, sx)
    pts <- tsx$pts
    xeval <- tsx$xeval
  }
  return(list(sample = s, xeval = xeval, pts = pts))
}

```

Problem 3

We will explore the need in importance sampling that the sampling density have heavier tails than the density of interest. We will estimate EX and $E(X^2)$ with respect to density f in this problem.

(a)

Suppose $f \sim N(0, 1)$ with sampling density $g \sim t(df = 3)$. We sample $m = 10000$ points to extract histograms of estimates and weights in order to get an idea whether $Var(\hat{\mu})$ is large.

As you can see from the results shown below, the variance is not so large since there are not many extreme weights in the samples.



```
## (a) we want to estimate EX and VX for a standrad normal drawn from a t
## with 3 degrees of freedom
m <- 10000 # number of samples for each estimator
set.seed(0)
x <- rt(m, df = 3) # i.e sample from g being a t with df=3
f <- dnorm(x) # density of x under f
g <- dt(x, df = 3) # density of x under g
w <- f/g # weights
mean(x * w) #EX

## [1] 0.02107

mean((x * w)^2) #VX

## [1] 0.9225

par(mfrow = c(1, 2), cex = 10, cex.main = 1.2)
hist(w)
hist(x * w)
```



(b)

Suppose $f \sim t(df = 3)$ with sampling density $g \sim N(0, 1)$. We sample $m = 10000$ points to extract histograms of estimates and weights in order to get an idea whether $Var(\hat{\mu})$ is large.

As the sampling density $t(df = 3)$ has less heavier tails than the density of interest $N(0, 1)$, there are many extreme weights that have a very strong influence on the estimation variance, as shown below.

```
## (b) we want to estimate EX and EX^2 for a t with v=3 drawn from a
## standard normal
m <- 10000 # number of samples for each estimator
set.seed(0)
x <- rnorm(m) # i.e sample from g being a standard normal
f <- dt(x, df = 3) # density of x under f
g <- dnorm(x) # density of x under g
w <- f/g # weights
mean(x * w) #EX



## [1] 0.004417

mean((x * w)^2) #VX

## [1] 4.846

par(mfrow = c(1, 2), cex = 10, cex.main = 1.2)
hist(log(w))
hist(log(x * w))

## Warning: NaNs produced
```

Histogram of $\log(w)$ Histogram of $\log(x * w)$ 

Problem 4

This is a truncated/censored regression problem. In a given i.i.d. sample, stochastically, c of the n original observations will be censored (point y_i will be denoted as NA in the simulated test sample if and only if $y_i > \tau$). Without loss of generality, we can separate the indices set into $i = 1, \dots, c$ with NA and $i = c+1, \dots, n$ with y_i . Here, we assume a simple linear regression model, $Y \sim X\beta + \epsilon$ with $y_i \sim N(\beta_0 + \beta_1 x_i, \sigma^2)$, i.e. define $\beta = (\beta_0, \beta_1)$, $\mu_i = (1, x_i)^T \beta$

(a)

EM algorithm is designed to estimate the parameter set, $\theta = (\beta_0, \beta_1, \sigma^2)$. Here, we take the complete data to be the available data plus the actual values of the truncated observations, i.e. $\{Y, Z\} = \{\{y_i\}, \{z_i\}\}$, of which $z_i, i = 1, \dots, c$ are values in place of the initially NA censored observations.

and they are functions of $\theta^{(t)}$, constant in terms of the maximization step. So the expected complete log likelihood is

$$\log N(y_i; \theta) = -\log \sqrt{2\pi\sigma^2} - \frac{(y_i - \mu_i)^2}{2\sigma^2}; \mu_i = \beta_0 + \beta_1 x_i$$

$$\log f(Y, Z; \theta) = \sum_{i=c+1}^n \log N(y_i; \theta) + \sum_{i=1}^c \log N(z_i; \theta)$$

$$E[\log N(z_i; \theta) | Y, X, \theta^{(t)}] = E[\log N(z_i; \theta) | z_i > \tau, \theta = \theta^{(t)}] = -\log \sqrt{2\pi\sigma^2} - \frac{1}{2\sigma^2} E[(z_i - \mu_i)^2 | z_i > \tau, \theta = \theta^{(t)}]$$

Now, let $m_i^{(t)} = E(z_i | z_i > \tau, \theta = \theta^{(t)})$ and $v_i^{(t)} = V(z_i | z_i > \tau, \theta = \theta^{(t)})$, where $m_i^{(t)} = \mu_i^{(t)} + \sigma^{(t)} \rho(\tau_i^{*(t)})$ and $v_i^{(t)} = (\sigma^2)^{(t)} (1 + \tau_i^{*(t)} \rho(\tau_i^{*(t)}) - \rho^2(\tau_i^{*(t)}))$ of which $\mu_i^{(t)} = \beta_0^{(t)} + \beta_1^{(t)} x_i$, $\tau_i^{*(t)} = (\tau - \mu_i^{(t)})/\sigma^{(t)}$,

$$\rho(\tau_i^{*(t)}) = \frac{\phi(\tau_i^{*(t)})}{1 - \Phi(\tau_i^{*(t)})} \text{ and } \phi(\cdot) \text{ is the standard normal PDF, } \Phi(\cdot) \text{ is the standard normal CDF.}$$

We have $E[(z_i - \mu_i)^2 | z_i > \tau, \theta = \theta^{(t)}] = (m_i^{(t)} - \mu_i^{(t)})^2 + v_i^{(t)}$; So,

$$Q(\theta; \theta^{(t)}) = E[\log f(Y, Z; \theta) | Y, X, \theta^{(t)}] = \sum_{i=c+1}^n \log N(y_i; \theta) + \sum_{i=1}^c \log N(m_i^{(t)}; \theta) + \sum_{i=1}^c \frac{v_i^{(t)}}{-2\sigma^2}$$

Then, the maximization step would give us updates on θ very similar to standard linear regression;

$$\beta^{(t+1)} = (X^T X)^{-1} X^T W^{(t)}, \text{ where } W^{(t)} = \{Y, M^{(t)}\} = \{\{y_i\}, \{m_i^{(t)}\}\};$$

$$(\sigma^2)^{(t+1)} = \frac{1}{n} \left[\sum_{i=c+1}^n (y_i - \mu_i^{(t)})^2 + \sum_{i=1}^c (m_i^{(t)} - \mu_i^{(t)})^2 + \sum_{i=1}^c v_i^{(t)} \right]$$

(b)

The reasonable starting values for the 3 parameters as functions of the observations are:

$$\beta_0^{(0)} = \bar{y} - \beta_1^{(0)} \bar{x}$$

$$\beta_1^{(0)} = [(n-c)\bar{y} + c\tau]/\bar{x}$$

$$(\sigma^2)^{(0)} = \frac{1}{n-c} \left[\sum_{i=c+1}^n (y_i - \beta_0^{(0)} - \beta_1^{(0)} x_i)^2 \right]$$

(c)

In order to test the implemeted EM algorithm for parameter estimation, we have simulated data for this purpose. For a sample size of $n = 100$, we make the true parameters such that with complete data, $\hat{\beta}_1/se(\hat{\beta}_1) \approx 3$, i.e. $\sigma^2 = (\beta_1/3)^2 \sum_{i=1}^n (x_i - \bar{x})^2$. The R function *EM.censor* takes in x, y, τ and estimate θ

with initialization from (b) and $lm()$ for updating β . The criteria for deciding when to stop the optimization is with $diff < tol$, i.e. the difference between the last and current θ is below the fixed tolerance value.

```
#### (c) EM implementation
EM.censor <- function(x, y, tau) {
  obs <- which(!is.na(y)) # observed y indices
  cen <- which(is.na(y)) # censored y indices
  c <- length(cen)
  n <- length(y) # sample size
  ny <- y #censored y's to be updated by EM
  # simulation setup
  tol <- 1e-08
  iter <- 0
  lim <- 1000
  beta0.trace <- matrix(NA, lim, 1)
  beta1.trace <- matrix(NA, lim, 1)
  sigma2.trace <- matrix(NA, lim, 1)
  more = TRUE
  # initialization
  beta1 <- ((n - c) * mean(y[obs]) + c * tau)/(n * mean(x))
  beta0 <- mean(y[obs]) - beta1 * mean(x)
  sigma2 <- sum((y[obs] - beta0 - beta1 * x[obs])^2)/(n - c)
  # auxiliary function
  rho <- function(x) {
    return(dnorm(x)/(1 - pnorm(x)))
  }
  # main optimization
  while (more) {
    iter <- iter + 1
    # E-step
    mu <- beta0 + beta1 * x
    tauS <- (tau - mu)/sqrt(sigma2)
    ny[cen] <- mu[cen] + sqrt(sigma2) * rho(tauS[cen]) # E(Z)
    vy <- sigma2 * (1 + tauS[cen] * rho(tauS[cen]) - rho(tauS[cen])^2) # V(Z)
    # M-step
    fit <- lm(ny ~ x) # LS fit for beta MLEs
    b0 <- fit$coefficients[1]
    b1 <- fit$coefficients[2]
    sig2 <- (sum(fit$residuals^2) + sum(vy))/n # MLE of sigma2
    # convergence
    diff.th <- abs(beta0 - b0) + abs(beta1 - b1) + abs(log(sigma2) - log(sig2))
    more <- (diff.th > tol) # check tol
    # update theta
    beta0 <- b0
    beta1 <- b1
    sigma2 <- sig2
    # record trace
    beta0.trace[iter] <- beta0
    beta1.trace[iter] <- beta1
    sigma2.trace[iter] <- sigma2
  }
  # results
  theta.result <- c(beta0, beta1, sigma2)
  names(theta.result) <- c("beta0", "beta1", "sigma2")
}
```

```

print(theta) #global true parameter
print(iter)
print(theta.result)
# plots
par(mfrow = c(3, 1), mar = c(4, 4, 2, 2), cex = 10, cex.main = 1.2)
plot(1:iter, beta0.trace[1:iter], xlab = "iteration", ylab = "beta0", "l")
plot(1:iter, beta1.trace[1:iter], xlab = "iteration", ylab = "beta1", "l")
plot(1:iter, sigma2.trace[1:iter], xlab = "iteration", ylab = "sigma2",
     "l")
par(mfrow = c(1, 1), cex = 10, cex.main = 1.2)
plot(x, y, ylim = c(min(y[obs]) - 1, max(y[obs]) + 2)) # plot data
points(x[cen], ny[cen], type = "p", pch = 20)
lines(c(sort(x)[1], sort(x)[n]), c(beta0 + beta1 * sort(x)[1], beta0 + beta1 *
    sort(x)[n])) # LS line
}

```

The test case (a) with a modest proportion of exceedances expected is shown below.

```

#### test data generation
set.seed(0)
n <- 100 # sample size
c <- n * 0.2 # test (a)/(b) exceedance=20%/80%
# true parameter values
theta <- rep(NA, 3)
names(theta) <- c("beta0", "beta1", "sigma2")
theta["beta0"] <- 0.5
theta["beta1"] <- 2
x <- runif(n)
theta["sigma2"] <- (theta["beta1"]/3)^2 * sum((x - mean(x))^2)
# simulated data
e <- rnorm(n, mean = 0, sd = sqrt(theta["sigma2"]))
y <- theta["beta0"] + theta["beta1"] * x + e
sort.y <- sort(y, decreasing = TRUE)
tau <- sort.y[c + 1] # threshold
y[which(y > tau)] <- as.numeric("NA")


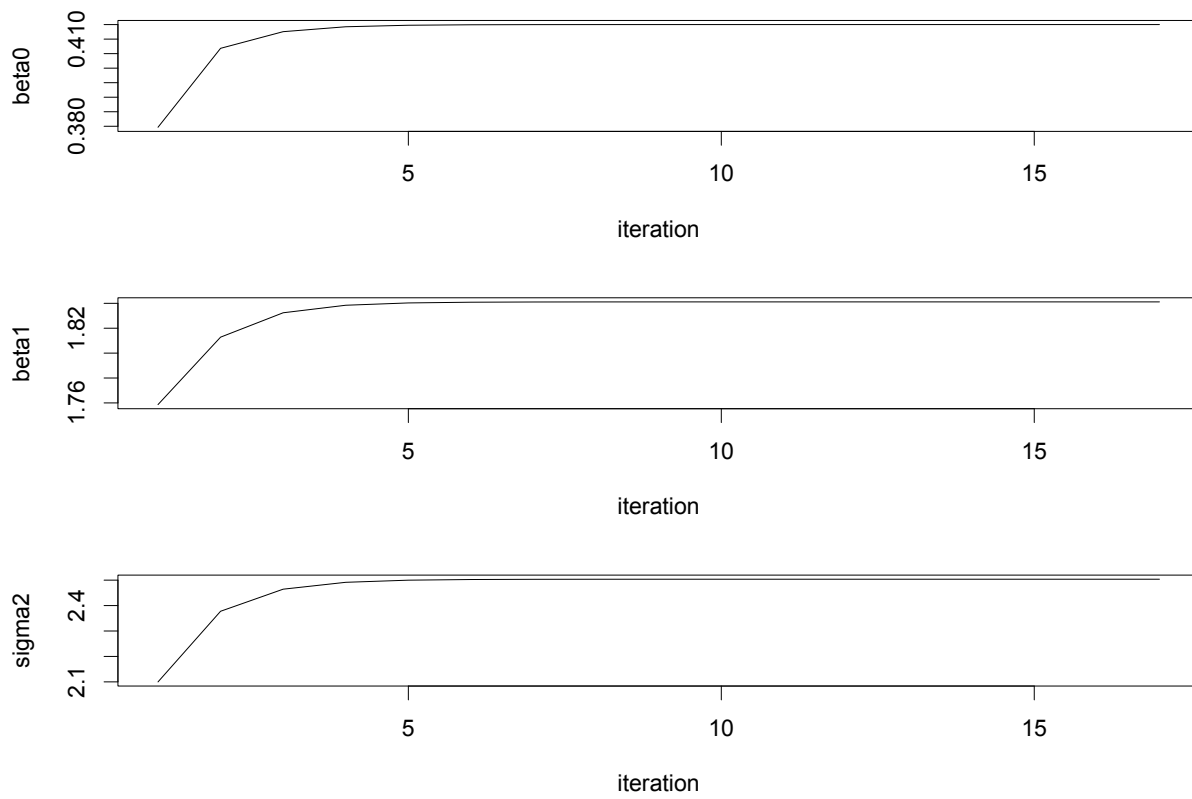
```

```

#### EM result
EM.censor(x, y, tau)

##  beta0  beta1 sigma2
##  0.500  2.000  3.211
## [1] 17
##  beta0  beta1 sigma2
##  0.415  1.841  2.504





```

The test case (b) with a high proportion of exceedances expected is shown below.

```
#### EM result
EM.censor(x, y, tau)
```

```
##  beta0  beta1  sigma2
##  0.500  2.000  3.211
##  [1] 216
##  beta0  beta1  sigma2
## -0.1705  1.8967  1.1354
```



(d)

A different approach to this problem will be just directly minimize the negative log likelihood of the data. For the censored observations, we just involve the likelihood terms, $P(Y_i > \tau), Y_i \sim N(\mu_i, \sigma^2)$. We use *optim()* with BFGS option in R to estimate the parameters and their standard errors.

```
#### (d) BFGS implementation
BFGS.censor <- function(x, y, tau) {
  obs <- which(!is.na(y)) # observed y indices
  cen <- which(is.na(y)) # censored y indices
  c <- length(cen)
  n <- length(y) # sample size
  # initialization
  beta1 <- ((n - c) * mean(y[obs]) + c * tau)/(n * mean(x))
  beta0 <- mean(y[obs]) - beta1 * mean(x)
  sigma2 <- sum((y[obs] - beta0 - beta1 * x[obs])^2)/(n - c)
  init0 <- c(beta0, beta1, sigma2)
  # log liklihood
  pp.lik <- function(par, x, y, tau, obs, cen) {
    b0 <- par[1]
    b1 <- par[2]
    sig2 <- par[3]
    mu <- b0 + b1 * x
    if (sig2 <= 0) {
      sig <- 1e-06
    } else {
      sig <- sqrt(sig2)
    }
    obs.lik <- sum(log(dnorm(y[obs], mean = mu[obs], sd = sig)))
    cen.lik <- sum(log(pnorm(tau, mean = mu[cen], sd = sig, lower.tail = F)))
    return(-(obs.lik + cen.lik))
  }
  # optimize
  opt <- optim(init0, pp.lik, x = x, y = y, tau = tau, obs = obs, cen = cen,
    method = "BFGS", control = list(trace = TRUE, parscale = c(1, 1, 10)),
    hessian = TRUE)
  # results
  print(theta)
  cat(opt$par, "\t theta \n")
  cat(sqrt(diag(solve(opt$hessian))), "\t se(theta) \n")
}
```

The test case (a) with a modest proportion of exceedances expected is shown below. It requires *iter* = 10 for BFGS, while *iter* = 17 for EM.

```
#### BFGS result
BFGS.censor(x, y, tau)
```

```
## initial value 191.328206
## iter 10 value 171.107012
## final value 171.049515
## converged
## beta0 beta1 sigma2
## 0.500 2.000 3.211
## 0.415 1.841 2.504 theta
## 0.3506 0.603 0.4165 se(theta)
```

The test case (b) with a high proportion of exceedances expected is shown below. It requires $iter = 42$ for BFGS, while $iter = 216$ for EM.

```
#### BFGS result
```

```
BFGS.censor(x, y, tau)
```

```
## initial value 187.280367
## iter 10 value 108.318570
## iter 20 value 98.301409
## iter 30 value 80.195175
## iter 40 value 58.781916
## final value 55.725382
## converged
## beta0 beta1 sigma2
## 0.500 2.000 3.211
## -0.1705 1.897 1.136 theta
## 0.3105 0.6849 0.4201 se(theta)
```

Problem 5

(a)

The data are measurements of flux from the supernova as a function of (log) wavelength and time. Our goal is to estimate flux (Y_{wt}) as a smooth 2-dimensional function of log wavelength (w) and time t pairs. We assume a normal likelihood as a function of parameters, $\theta = \{\kappa, \lambda, \sigma^2, \rho_w, \rho_t, \tau^2, \alpha\}$, $Y \sim N(\mu_\theta, C_\theta)$.



Here, μ_θ is a vector of value for observation

$E(Y_{wt}) = \mu(t; \theta) = \kappa f(\frac{t}{\lambda})$; and C_θ is a matrix with entries of pairwise covariances of the observations

$$\text{Cov}(Y_{wt}, Y_{w't'}) = \sigma^2 \exp\left(\frac{-|w - w'|}{\rho_w}\right) \exp\left(\frac{-|t - t'|}{\rho_t}\right) + \tau^2 I(t = t') + \alpha v_{wt}^2 I(w = w', t = t').$$

So the plots of data with respect to log wavelength and time are shown below.

```
#### 5. optim
rm(list = ls(all = TRUE)) # remove all objects
source("ps5prob5.R")
data$logw <- log(data$wavelength) # log wavelength
data$verr <- (data$fluxerror)^2 # square fluxerror
## (a) plot flux data
attach(data)
par(mfrow = c(1, 2), cex = 10, cex.main = 1)
plot(logw, flux, xlab = "log(wavelength)", ylab = "flux")
plot(time, flux, xlab = "time", ylab = "flux")
```



```
detach(data)
```

(b)

A reasonable set of initial values based on the scale of the flux values and the log wavelength and time values are given below inline with the code.

```
## (b) initialize
theta <- rep(NA, 7)
names(theta) <- c("kappa", "lambda", "sigma2", "rho_w", "rho_t", "tau2", "alpha")
theta["lambda"] <- 0.8
theta["kappa"] <- mean(data$flux)/mean(meanpts$value)
Ef <- meanfunc(theta, data$time)
resf <- data$flux - Ef
Covf <- outer(resf, resf)
theta["sigma2"] <- mean(diag(Covf)) * 0.3
theta["rho_w"] <- diff(range(data$logw))
theta["rho_t"] <- diff(range(data$time))
theta["tau2"] <- mean(diag(Covf)) * 0.3
theta["alpha"] <- mean(diag(Covf))/mean(sqrt(data$verr))
print(theta)

##      kappa      lambda      sigma2      rho_w      rho_t      tau2      alpha
## 0.335992 0.800000 0.003243 0.025623 38.916000 0.003243 42.993455
```

(c)

We use *optim()* to optimize the log likelihood to find estimates and standard error for θ . The results below showed that we have tried 2 different methods for finding the global optimum, the default Nelder-Mead and the BFGS. Also, *parscale* is used in the optimization to make the optimization reliable. However, due to the complexity in the likelihood function, we relaxed the convergence a little to reduce the optimization time.

```
## (c) optimize
covfunc <- with(list(verrs = data$verr), function(theta, wavel, times) {
  n <- length(times)
  CV <- matrix(0, n, n)
  scale_w <- wavel/theta["rho_w"]
  scale_t <- times/theta["rho_t"]
  alpv2 <- theta["alpha"] * verrs
  for (i in 1:n) {
    for (j in 1:n) {
      del_w <- abs(scale_w[i] - scale_w[j])
      del_t <- abs(scale_t[i] - scale_t[j])
      cv <- theta["sigma2"] * exp(-del_w) * exp(-del_t)
      if (!del_t) {
        cv <- cv + theta["tau2"]
        if (!del_w) {
          cv <- cv + alpv2[i]
        }
      }
      CV[i, j] <- cv
    }
  }
  CV
})
```

```

# negative log likelihood
require(mvtnorm)
nn.lik <- with(list(f = data$flux, w = data$logw, t = data$time), function(par) {
  MU <- meanfunc(par, t)
  CV <- covfunc(par, w, t)
  ll <- sum(sum(dmvnorm(f, MU, CV, log = T)))
  -ll
})

init0 <- theta
opt1 <- optim(par = init0, fn = nn.lik, control = list(trace = TRUE, parscale = init0,
  reltol = 1e-04, maxit = 80), hessian = TRUE)

## Nelder-Mead direct search function minimizer
## function value for initial parameters = -1610.819434
## Scaled convergence tolerance is 0.161082
## Stepsize computed as 0.100000
## BUILD      8 -1597.511663 -1623.889530
## EXTENSION  10 -1607.195503 -1647.098990
## LO-REDUCTION 12 -1608.778113 -1647.098990
## LO-REDUCTION 14 -1610.596435 -1647.098990
## LO-REDUCTION 16 -1610.819434 -1647.098990
## EXTENSION  18 -1611.126717 -1666.114092
## LO-REDUCTION 20 -1621.006730 -1666.114092
## EXTENSION  22 -1623.889530 -1690.194185
## LO-REDUCTION 24 -1631.267141 -1690.194185
## LO-REDUCTION 26 -1635.133629 -1690.194185
## EXTENSION  28 -1640.554216 -1729.373748
## LO-REDUCTION 30 -1647.098990 -1729.373748
## EXTENSION  32 -1663.290820 -1761.646704
## LO-REDUCTION 34 -1666.114092 -1761.646704
## LO-REDUCTION 36 -1684.387736 -1761.646704
## LO-REDUCTION 38 -1689.328919 -1761.646704
## LO-REDUCTION 40 -1690.194185 -1761.646704
## HI-REDUCTION 42 -1722.380543 -1761.646704
## LO-REDUCTION 44 -1723.828526 -1761.646704
## HI-REDUCTION 46 -1729.373748 -1761.646704
## LO-REDUCTION 48 -1741.130855 -1761.646704
## LO-REDUCTION 50 -1748.173626 -1761.646704
## LO-REDUCTION 52 -1749.897789 -1761.646704
## LO-REDUCTION 54 -1752.170165 -1761.646704
## EXTENSION  56 -1756.418165 -1769.176145
## LO-REDUCTION 58 -1756.762710 -1769.176145
## LO-REDUCTION 60 -1757.052092 -1769.176145
## EXTENSION  62 -1759.142872 -1778.579957
## LO-REDUCTION 64 -1759.596655 -1778.579957
## LO-REDUCTION 66 -1760.845131 -1778.579957
## LO-REDUCTION 68 -1761.646704 -1778.579957
## EXTENSION  70 -1763.350160 -1792.370206
## LO-REDUCTION 72 -1768.382542 -1792.370206
## LO-REDUCTION 74 -1769.176145 -1792.370206
## LO-REDUCTION 76 -1769.458449 -1792.370206
## EXTENSION  78 -1771.363143 -1814.308692

```

```
## LO-REDUCTION      80 -1775.572363 -1814.308692
## Exiting from Nelder Mead minimizer
##      82 function evaluations used

cat("theta \n", opt1$par, "\n")

## theta
##  0.261 1.21 0.001584 0.04837 60.05 0.003646 16.87

cat("se(theta) \n", sqrt(diag(solve(opt1$hessian))), "\n")

## se(theta)
##  0.04703 0.3371 0.0001027 0.007372 11 NaN NaN
```

```
init0 <- theta
opt2 <- optim(par = init0, fn = nn.lik, method = "BFGS", control = list(trace = TRUE,
  parscale = init0, reltol = 1e-04, maxit = 20))

## initial value -1610.819434
## iter 10 value -1929.347625
## final value -1941.131307
## converged

cat("theta \n", opt2$par, "\n")

## theta
##  0.1759 0.8808 0.0005201 0.04311 30.64 0.00256 -1.27
```


Problem 6

In the work that won the Netflix Prize done by C. Volinsky etc. at AT&T labs, a variation on the singular value decomposition (SVD) of the matrix factorization methodology is employed.

Given a matrix R that represents m users (individual Netflix members) and n items (movies) with entries of movie ratings r_{ui} , $R = \{r_{ui}\}_{1 \leq u \leq m, 1 \leq i \leq n}$, the SVD factorization-based methodology computes the best rank- f approximation $R^f = P_{m \times f} Q_{n \times f}^T$, where $f \leq m, n$.

Due to the fact that most entries of R in this movie recommendation problem are unknown, the SVD here is used to extend the given data by filling in the unknown ratings by estimation $r_{ui} \sim R_{ui}^f = p_u^T \cdot q_i$, where p_u is the u -th row of P corresponding to user u and q_i is the i -th row of Q corresponding to item i .

An EM-based algorithm is employed to iteratively compute the R matrix SVD by doing least square minimization of $\|R - PQ^T\|_F$ when alternating the fixed point between matrix P and Q , i.e.

$$Q^T \leftarrow (P^T P)^{-1} P^T R$$

$$P \leftarrow RQ(Q^T Q)^{-1}$$

Shrinkage is integrated to alleviate the overfitting problem and further reduce the estimation error,

$$Err(P, Q) \equiv \sum_{(u,i)} (r_{ui} - p_u^T q_i)^2.$$

To be more specific, assuming that we have already computed the first $f - 1$ columns of matrices P, Q . The pseudo code for computing the f -th column of matrices P, Q is given below [1]:

```
# Problem 5.6: EM based SVD variant
f <- f + 1 # iteration
computeNext <- function(r, P, Q) {
  #known ratings, user factors, item factors
  # compute the f-th column of P,Q to fit given ratings columns 1,...,f-1
  # were already computed
  a <- alpha
  e <- eps
  # compute residuals portion not explained by previous factors
  for (u in 1:m) {
    for (i in 1:n) {
      r[u, i] <- r[u, i] - crossprod(P[u, 1:(f - 1)], Q[i, 1:(f - 1)])
      n <- s[u, i] # support behind r
      r[u, i] <- n * r[u, i] / (n + a * f) #shrinkage
    }
  }
  # compute the f-th factor for each user and item by solving many least
  # square problems, each with a single unknown
  while (err(P, Q) < 1 - e) {
    for (u in 1:m) {
      P[u, f] <- crossprod(r[u, ], Q[, f]) / crossprod(Q[, f])
    }
    for (i in 1:n) {
      Q[i, f] <- crossprod(r[, i], P[, f]) / crossprod(P[, f])
    }
  }
  return(list(P = P, Q = Q))
}
```

At the end of the process, we obtain an approximation of all ratings in the form of a matrix product PQ^T , with predictions for unrated movies. The eigenvectors from SVD are user and item factors, while the eigenvalues are confidence indicators for the predictions.

References

- [1] R. Bell, Y. Koren, C. Volinsky, "Modeling relationships at multiple scales to improve accuracy of large recommender systems", *Proceedings of the 13th ACM SIGKDD International Conference on Knowledge Discovery and Data mining*, pp. 95-104, Aug. 2007, San Jose, CA.